



# SLIIT

*Discover Your Future*

# IT2060/IE2061

## Operating Systems and System Administration

### Lecture 06

## Introduction to Deadlock

**U. U. Samantha Rajapaksha**

M.Sc.in IT, B.Sc.(Engineering) University of Moratuwa

Senior Lecturer SLIIT

[Samantha.r@slit.lk](mailto:Samantha.r@slit.lk)



**SLIIT**  
**FACULTY OF COMPUTING**

# DEADLOCKS

- Several processes may compete for a finite number of resources, and some of them may wait for the resources forever because the resources are held by other waiting processes → deadlock.
- A set of processes is in a deadlock state if every process in the set is waiting for an event that can be caused only by another process in the set.

## Example

- System has two tape drives.
- P1 and P2 each hold one tape drive and each needs another one.

## Example

- Semaphores A and B, initialized to 1.

P0

P1

*wait (A)*

*wait (B)*

*wait (B)*

*wait (A)*

# System Model

- Resources are partitioned into several types, each consists of some number of identical *instances*.
  - **Identical**: allocation of *any* instance of the type will satisfy process's request.
  - Resources may be physical resources (printers, tape drives, CPU cycles), or logical resources (files, semaphores, and monitors).
  - A **pre-emptible** resource is one that can be taken away from a process with no ill effect to the process; e.g., memory.
  - A **non-preemptible** resource is one that cannot be taken away from its user since it will make the user fails; e.g., printers
    - In general, potential deadlocks involve this resource type.
- Each process uses a resource as follows:
  - **Request** the resource; a process must wait if the resource is being used by another process.
  - **Use** the resource; e.g., the process can print on the printer.
  - **Release** the resource.

# Necessary conditions for deadlock

Four conditions must hold for a deadlock to occur (Coffman et al.):

1. **Mutual exclusion condition.** Only one process at a time can use the resource.

- **or** each resource is either currently assigned to exactly one process or is available.

2. **Hold and wait condition.** A process holding at least one resource is waiting to acquire additional resources held by other processes.

3. **No pre-emption condition.** A resource can be released only voluntarily by the process holding it after that process has completed its task.

4. **Circular wait condition.** There exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .

**Note:** the four conditions are not completely independent, e.g., the circular-wait condition implies the hold-and-wait condition.



# Deadlock Modelling

- Deadlocks can be described more precisely in terms of a directed graph  $G(V, E)$ 
  - called *System resource-allocation graph*
- $V$  is partitioned into two types:
  - Set of processes in the system:  $P = \{P_1, P_2, \dots, P_n\}$ .
  - Set of all resource types in the system:  $R = \{R_1, R_2, \dots, R_n\}$
- *Request edge* – directed edge  $P_i \rightarrow R_j$ 
  - process  $P_i$  requests an instance of resource  $R_j$
- *Assignment edge* – directed edge  $R_j \rightarrow P_i$ 
  - an instance of resource  $R_j$  has been allocated to process  $P_i$


# Model Symbols

- Process: ○


Resource type with 4 instances:



- $P_i$  requests  $R_j$ :




$P_i$  uses  $R_j$ :



## Example: Deadlock


Printer



Tape drive

- \* If the graph contains **no cycles**, no process in the system is deadlocked.
- \* If the graph contains **a cycle**, deadlock *may* exist.
  - If each resource type has **one instance**, **cycle means deadlock**.
  - If each resource type has **several instances**, cycle is necessary but **not sufficient condition for deadlock**.

## Example: resource allocation graph (with no cycles)



### The sets P, R, and E:

$$P = \{P_1, P_2, P_3\}$$

$$R = \{R_1, R_2, R_3, R_4\}$$

$$E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, R_1 \rightarrow P_2, R_2 \rightarrow P_2, R_2 \rightarrow P_1, R_3 \rightarrow P_3\}$$

### Resource instances:


- One instance of resource type  $R_1$
- Two instances of resource type  $R_2$
- One instance of resource type  $R_3$
- Three instances of resource type  $R_4$

### Process states:

- $P_1$  is holding an instance of  $R_2$ , and waiting for an instance of  $R_1$
- $P_2$  is holding an instance of  $R_1$  and  $R_2$ , and is waiting for an instance of  $R_3$
- $P_3$  is holding an instance of  $R_3$


ets  
the

## A cycle and deadlock



## A cycle but no deadlock

P<sub>4</sub> can release R<sub>2</sub> which gets allocated to P<sub>3</sub>; breaking the cycle





## Three Methods for handling deadlock

- Use a protocol to ensure that the system will *never* reaches deadlock
  - Using *deadlock prevention* and/or *deadlock avoidance* techniques
- Allow the system to enter a deadlock state and then recover
  - needs *deadlock detection* and *deadlock recovery* algorithms
- Ignore the problem and pretend that deadlocks never occur in the system
  - used by most OS's, including UNIX
  - Also called the **ostrich** algorithm!

# Deadlock prevention

- Restrain the ways resource requests can be made
  - Use a set of methods to ensure that **any one** of the four deadlock conditions cannot hold

## (1) Deny mutual exclusion

- Not required for sharable resources (e.g., read-only files, cannot be in deadlock)
- Must hold for non-sharable resources (a printer cannot be simultaneously shared by several processes)
- **In general**, it is not possible to prevent deadlock by denying mutual-exclusion condition since some resources are non-sharable

## (2) Deny hold and wait

- Must guarantee that whenever a process requests a resource, it does not hold any other resources

### Options:

- Each process is granted all resources before it starts
- Allows a process to request resources only when it has none
  - If a process needs more resources, release all resources before requesting new ones

### Problem:

- Resource utilisation is low
- Possible starvation.
  - A process that needs popular resources may have to wait indefinitely

## Deadlock prevention (cont.)

### (3) Prevent no pre-emption (i.e., allow pre-emption)

- When a process holding some resources requests other resource that cannot be immediately allocated, it must release all resources currently being held
  - The pre-empted resources are added to the process's list of requested resources
  - The process is restarted when it regains its old resources and obtains the new one it is requesting

#### Problem:

- Can be applied easily to resources whose state can be saved easily (e.g., memory), but not so easily for others (e.g., printer)

## Deadlock prevention (cont.)

### (4) Deny circular wait

- All resource types are ordered, e.g.,
  - $F(\text{card reader}) = 1$                        $F(\text{disk drive}) = 5$
  - $F(\text{tape drive}) = 7$                        $F(\text{printer}) = 12$
- Each process must request increasing order of resources
- Protocol:
  - Each process requests resources in increasing order
  - Initially a process can request for any  $R_i$
  - After that, it can request  $R_j$  only if  $F(R_j) > F(R_i)$
- **Problem:** It may be impossible to find a resource ordering that satisfies everyone



# Deadlock avoidance

- The system must have some additional *a priori* information about which resources a process will request and use during its lifetime
  - With the additional information, the system can decide for each request whether or not the process should wait
  - The simplest and most useful model requires that each process declare the *maximum* number of resources of each type that it may need
- The deadlock-avoidance algorithm dynamically examines the resource-allocation *state* to ensure that **there can never be a circular-wait condition**
- A resource-allocation *state* is defined by:
  - The number of **available** and **allocated** resources, and
  - The **maximum demands** of the processes


# Safe State

- When a process requests an available resource, the system checks if its allocation keeps the system in. *safe state*
- The system is in *safe state* if there exists a *safe sequence* of all processes
- A sequence  $\langle P_1, P_2, \dots, P_n \rangle$  is *safe* if, for each  $P_i$ , the resources requested by  $P_i$  can be allocated from the currently available resources + resources held by all  $P_j$ , with  $j < i$ 
  - If  $P_i$ 's resource needs are not immediately available,  $P_i$  waits until all  $P_j$  have finished
  - When all  $P_j$  are finished,  $P_i$  obtains the needed resources, executes, returns the allocated resources, and terminates
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on

## Safe State (cont.)

### Basic facts

- If a system is in safe state → no deadlocks
- If a system is in unsafe state → possibility of deadlock
- Avoidance ensures that the system never enters an unsafe state
- A process requesting for a currently available resource may have to wait
  - Thus, resource allocation is lower than without deadlock avoidance algorithm



## Example

Consider a system with 12 resources of the same type, and 3 processes with the following resource needs and allocation

	<u>Maximum needs</u>	<u>Allocation</u>	<u>Current need</u>
$P_0$	10	5	5
$P_1$	4	2	2
$P_2$	9	2	7

- At time  $t_0$ , **available resource = 3**, and the system is in safe state
  - There is a safe sequence  $\langle P_1, P_0, P_2 \rangle$
- What if at  $t_1$  one more resource is allocated to process  $P_2$ ?
  - The system is in unsafe state
    - Deadlock can occur





# Resource-Allocation Graph Algorithm

- *Claim edge*  $P_i \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$ 
  - represented by a dashed line
- *Claim edge* converts to *request edge* when a process requests a resource
- When a resource is released by a process, *assignment edge* converts to a *claim edge*
- Resources must be claimed *a priori* in the system
- Need a cycle detection algorithm  $\rightarrow O(n^2)$
- This algorithm **can not be used** for system comprising resource types with **multiple instances**

## Example

Suppose  $P_2$  requests  $R_2$

Although  $R_2$  is currently free, allocating it to  $P_2$  may lead to unsafe state (a cycle in right figure)



# Banker's Algorithm

- The algorithm for a system comprising resource types with **multiple instances**
- Similar to a bank: never allocates its available cash if it can no longer satisfy the needs of all customers
- Each process must *a priori* claim maximum number of instances of each resource type that it may need
- When a process requests a resource:
  - It may have to wait (if resource allocation may lead to unsafe state) until some other process releases enough resources
- When a process gets all its resources:
  - It must return them in a finite amount of time

# Banker's Algorithm (cont.)

## Algorithm

Let  $n$  = number of processes, and  $m$  = number of resource types

### Data structures:

- *Available*: Vector of length  $m$ 
  - $available[j] = k$ ; means  $k$  instances of resource type  $R_j$  are available
- *Max*:  $n \times m$  matrix
  - $Max[i, j] = k$ ; means process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
- *Allocation*:  $n \times m$  matrix
  - $Allocation[i, j] = k$ ; means process  $P_i$  is currently allocated  $k$  instances of resource type  $R_j$
- *Need*:  $n \times m$  matrix
  - $Need[i, j] = k$ ; means process  $P_i$  may need  $k$  more instances of resource type  $R_j$  to complete its task.
  - $Need[i, j] = max[i, j] - allocation[i, j]$



# Implementation of the safety algorithm

// Time complexity =  $O(mn^2)$

1. Let *work* and *finish* be vectors of length *m* and *n*, respectively

initialise:

*work* = *available*

*finish* [*i*] = *false*      for *i* = 1, 2, ..., *n*

// Find an unfinished process *i*; it still needs resources

2. Find a value of *i* such that both:

- *finish*[*i*] = *false*, and
- *need*<sub>*i*</sub> ≤ *work*
- If no such *i* exists, go to step 4

// process *i* pretends to finish, so it releases its resources i.e., *allocation*<sub>*i*</sub>

3. *work* = *work* + *allocation*<sub>*i*</sub>

*finish*[*i*] = *true*

go to step 2

4. If *finish*[*i*] = *true* for all *i*, the system is in safe state.

## Resource-request algorithm for process $P_i$

$Request_i$  = request vector for process  $P_i$

If  $Request_i[j] = k$ , then process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If  $request_i \leq need_i$ , go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If  $request_i \leq available$ , go to step 3. Otherwise  $P_i$  must wait, since resources are not available.
3. The system pretends to allocate requested resources to  $P_i$  by modifying the state as follows:
  - $available = available - request_i$
  - $allocation_i = allocation_i + request_i$
  - $need_i = need_i - request_i$
  - If resulting state is safe, resources are allocated to  $P_i$
  - else  $P_i$  must wait, and the old resource-allocation state is restored.

# Example of Banker's algorithm

- 5 processes  $P_0$  through  $P_4$ ; 3 resource types A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time  $T_0$ :

	<i>Allocation</i>			<i>Max</i>			<i>Available</i>			<i>Need</i>		
	A	B	C	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	5	3	3	3	2	7	4	3
$P_1$	2	0	0	3	2	2				1	2	2
$P_2$	3	0	2	9	0	2				6	0	0
$P_3$	2	1	1	2	2	2				0	1	1
$P_4$	0	0	2	4	3	3				4	3	1

- ★ The content of matrix *Need* is defined to be *Max* – *Allocation*
- ★ The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies the safety criteria

## Example ( $P_1$ requests (1,0,2)):

- Check that  $request \leq need$  (that is,  $(1, 0, 2) \leq (1, 2, 2) \rightarrow$  true
- Check that  $request \leq available$  (that is,  $(1, 0, 2) \leq (3, 3, 2) \rightarrow$  true

### Before Adjustment

	Allocation		
	A	B	C
$P_0$	0	1	0
$P_1$	2	0	0
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

### After Adjustment

	Alloc.			Need			Avail.		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	4	3	2	3	0
$P_1$	3	0	2	0	2	0			
$P_2$	3	0	2	6	0	0			
$P_3$	2	1	1	0	1	1			
$P_4$	0	0	2	4	3	1			

\*  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  or  $\langle P_1, P_4, P_3, P_0, P_2 \rangle$  satisfies safety requirement

\* Can request for (3, 3, 0) by  $P_4$  be granted? (0, 2, 0) by  $P_0$ ?



# Deadlock detection


- If a system does not employ either a deadlock-prevention or a deadlock-avoidance algorithm, then a deadlock situation may occur
- Need a *deadlock detection* algorithm that examines the state of the system to determine whether a deadlock has occurred
- Need a *recovery* algorithm to recover from deadlock

## Deadlock detection for single instance of each resource type


- Maintain a *wait-for graph*
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- Periodically invoke an algorithm that searches for a cycle in the graph
  - An algorithm to detect a cycle in a graph requires  $O(n^2)$  operations,
    - $n$  is the number of vertices in the graph

# Example

## Resource Allocation Graph



## Wait for Graph



# Deadlock recovery

## 1) Terminate processes

- Kill (abort) all deadlocked processes
- Kill one process at a time until deadlock cycle eliminated
- In which order should we choose process to abort?
  - The process with lowest priority
  - How long the process has computed, and how much longer to completion
  - Resources the process has used
  - Resources the process needs to complete
  - How many processes will need to be terminated
  - Is process interactive or batch?

**Problem:** what if the process is in the middle of updating a file?

Aborting the process may lead to incorrect file

# Deadlock recovery

## 2) Pre-empt a resource from a process.

- How to select a victim (process) to minimize cost?
- Roll back the process to some safe state and restart from there
  - How do we find a safe state?
    - Easiest way: destroy the process and restart
    - Use checkpoints during execution
- Starvation – same process may always be picked as victim
  - How do we ensure no starvation?
    - Include number of rollbacks in cost factor