



SLIIT

Discover Your Future

IT2060/IE2061

Operating Systems and System Administration

Lecture 07

Process Synchronization

U. U. Samantha Rajapaksha

M.Sc.in IT, B.Sc.(Engineering) University of Moratuwa

Senior Lecturer SLIIT

Samantha.r@slit.lk



SLIIT
FACULTY OF COMPUTING

Chapter 5: Process Synchronization

- Background
- The Critical-Section Problem
- Semaphores
- Classic Problems of Synchronization
- Monitors

Background

- Processes can execute concurrently
 - May be interrupted at any time, partially completing execution
- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Illustration of the problem:

Suppose that we wanted to provide a solution to the consumer-producer problem that fills *all* the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers. Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.



Producer

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```



Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```


Race Condition

- **counter++** could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- **counter--** could be implemented as


```
register2 = counter
register2 = register2 - 1
counter = register2
```

- Consider this execution interleaving with “count = 5” initially:

S0: producer execute	<code>register1 = counter</code>	{register1 = 5}
S1: producer execute	<code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute	<code>register2 = counter</code>	{register2 = 5}
S3: consumer execute	<code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute	<code>counter = register1</code>	{counter = 6}
S5: consumer execute	<code>counter = register2</code>	{counter = 4}

Race condition

- *Race condition* is a situation where several processes access and manipulate the same data concurrently.
 - The outcome of the execution depends on particular order in which the access takes place.
- In order to prevent race condition on *counter*, we need to ensure that only one process at a time can be manipulating *counter*
 - We need some form of *process synchronization*.



Critical Section Problem

- Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code
 - Process may be changing common variables, updating table, writing file, etc
 - When one process in critical section, no other may be in its critical section
- **Critical section problem** is to design protocol to solve this
- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**

Critical Section

- General structu

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted

Solution (cont.)

Simplest solution:

Each process *disables* all interrupts just *after entering* its critical section and *re-enables* them just *before leaving* it.

- Not wise, because enabling/disabling interrupt is a privileged instruction.

Solution in kernel mode:

- **Preemptive kernel:** a process can be preempted while running in the kernel mode
 - Otherwise, the kernel is **nonpreemptive kernel:** allows the process to run until it exits kernel mode, blocks, or voluntarily yields CPU.
- Nonpreemptive kernel is free from race conditions on kernel data structure
 - However preemptive kernel is more responsive and suitable for real time system.

Solution for two processes, P_0 and P_1

Software-based

ALGORITHM 1

var *turn*: (0 .. 1); // Initially *turn* = 0; *turn* = *i* means P_i can enter its CS

Process P_i

repeat

while *turn* $\neq i$ **do** *no-op*;

 Critical Section

turn = *j*;

 Remainder Section

until *false*;

Process P_j

repeat

while *turn* $\neq j$ **do** *no-op*;

 Critical Section

turn = *i*;

 Remainder Section

until *false*;

- Satisfies mutual exclusion, but not progress requirement.
 - If *turn* = 0, P_1 cannot enter its CS even though P_0 is in its RS.
 - Taking turn is not good when one process is slower than other.
- *Busy waiting*: continuously testing a variable waiting for some value to appear
 - not good since it wastes CPU time

Solution for two processes (cont.)

Software-based

ALGORITHM 2

// Initially $flag[0] = flag[1] = false$; $flag[i] = true$ means P_i wants to enter its CS

var $flag$: **array** $[0 .. 1]$ **of** *boolean*;

Process P_i

repeat

$flag[i] = true$;

while $flag[j]$ **do** *no-op*;

Critical Section;

$flag[i] = false$;

Remainder Section;

until *false*;

Process P_j

repeat

$flag[j] = true$;

while $flag[i]$ **do** *no-op*;

Critical Section;

$flag[j] = false$;

Remainder Section;

until *false*;

- Satisfy mutual exclusion, but violates the progress requirement:

T_0 : P_0 sets $flag[0] = true$.

T_1 : P_1 sets $flag[1] = true$.

→ P_0 and P_1 are looping in their respective **while**.

Solution for two processes (cont.)

Software-based

// **Peterson's solution:** Combine shared variables of Algorithms 1 and 2

Process P_i

repeat

flag[i] = true;

turn = j;

while (*flag[j] and turn = j*) **do no-op;**

critical section

flag[i] = false;

remainder section

until false;

Process P_j

repeat

flag[j] = true;

turn = i;

while (*flag[i] and turn = i*) **do no-**

critical section

flag[j] = false;

remainder section

until false;

- Solves the critical-section problem for two processes.
 - **It meets all the three requirements**
- Proof: need to show that:
 - Mutual exclusion is preserved.
 - The progress requirement is satisfied.
 - The bounded waiting time requirement is met.
- For detailed proof, read the textbook.

Bakery Algorithm

- The solution to the critical section problem for n processes by Leslie Lamport
- Before entering its critical section, each process receives a number.
 - The holder of the smallest number enters the critical section.
 - If processes P_i and P_j receive the same number, if $i < j$, then P_i is served first; else P_j is served first.
- Notation (**ticket#**, process **id#**)
 - $(a, b) < (c, d)$ if $a < c$ or if $a = c$ and $b < d$.
 - $\max(a_0, \dots, a_{n-1})$ is a number k such that $k \geq a_i$ for $i = 0, \dots, n-1$.
- shared data
 - **var** *choosing*: **array**[0.. $n-1$] of *boolean*;
 - *number*: **array** [0 .. $n-1$] of *integer*;
- data structures are initialised to *false* and 0, respectively

Bakery Algorithm (contd.)

Process P_i

repeat

choosing[i] = *true*;

number[i] = $\max(\text{number}[0], \text{number}[1], \dots, \text{number}[n-1]) + 1$;

choosing[i] = *false*;

for $j = 0$ **to** $n-1$ **do**

begin

while *choosing*[j] **do** *no-op*;

while *number*[j] $\neq 0$ **and** (*number*[j], j) < (*number*[i], i) **do** *no-op*;

end;

critical section

number[i] = 0;

remainder section

until *false*

Synchronization Hardware

- There is no guarantee that the software-based solution will work correctly in all computer architectures.
- The simple solution to critical section problem: disable interrupt while a shared variable is being modified.
 - This solution is not feasible in multiprocessor. Why?
- Use special hardware instructions such as *Test-and-Set* and *Swap*.
 - *Test-and-set* or *Swap* is an atomic instruction: it can not be interrupted until it completes its execution

// Test and set the content of a word *atomically*

```
function Test-and-Set (var
    boolean: target)
begin
    Test-and-Set = target;
    target = true;
end;
```

// Swapping instruction is done *atomically*

```
procedure Swap (var
    boolean: a, b)
var boolean: temp;
begin
    temp = a;
    a = b;
    b = temp;
end;
```

How to use them?

Mutual Exclusion with Test-and-Set

var *boolean: lock*; *lock* is a shared variable, initially set to *false*.

```
Repeat // Process Pi
    while Test-and-Set (lock) do no-op;
        Critical Section
    lock = false;
        Remainder Section
until false;
```

Mutual Exclusion with Swap

```
Repeat // Process Pi
    key = true;
    repeat
        Swap (lock, key);
    until key = false;
        Critical section
    lock = false;
        Remainder section
until false;
```

```
Repeat // Process Pj
    while Test-and-Set (lock) do no-op;
        Critical Section
    lock = false;
        Remainder Section
until false;
```

```
Repeat // Process Pj
    key = true;
    repeat
        Swap (lock, key);
    until key = false;
        Critical section
    lock = false;
        Remainder section
until false;
```

- Both do not satisfy the bounded waiting requirement.

Correct solution with Test-and-set

Shared data: **var** *waiting*: **array**[0..*n*-1] **of** *boolean*; *lock*: *boolean*; //All initialized to false

Process P_i

var *j*: 0..*n*-1; *key*: *boolean*;

repeat

waiting[*i*] = *true*;

key = *true*;

while *waiting* [*i*] **and** *key* **do** // enter CS if either *waiting*[*i*] or *key* is false

key = Test-and-Set (*lock*); // *key* is false if *lock* is false

waiting[*i*] = *false*;

Critical Section

j = *i*+1 **mod** *n*;

while (*j* ≠ *i*) **and** **not** *waiting*[*j*] **do** // check if any P_j is waiting for CS

j = *j*+1 **mod** *n*

if *j* = *i* **then**

lock = *false*; // no other process is waiting for CS

else

waiting[*j*] = *false*; // P_j is waiting, let it enter CS next

Remainder Section

until *false*;

Proof: Read textbook.



Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for processes to synchronize their activities.
- Semaphore S – integer variable
- Can only be accessed via two indivisible (atomic) operations

- **wait()** and **signal()**

- Originally called **P()** and **V()**

- Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

- Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```

Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
 - Same as a **mutex lock**
- Can solve various synchronization problems
- Consider P_1 and P_2 that require S_1 to happen before S_2

Create a semaphore “**synch**” initialized to 0

P1:

```
S1;  
signal(synch);
```

P2:

```
wait(synch);  
S2;
```

- Can implement a counting semaphore S as a binary semaphore

Semaphore Implementation

- Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore at the same time
- Thus, the implementation becomes the critical section problem where the `wait` and `signal` code are placed in the critical section
 - Could now have **busy waiting** in critical section implementation
 - But implementation code is short
 - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution



Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue

```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore;
```


Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

Classical Problems of Synchroni

- Classical problems used to test new proposed synchronization schemes
 - Bounded-Buffer Problem
 - Readers and Writers Problem
 - Dining-Philosophers Problem

Bounded-Buffer Problem

- n buffers, each can hold one item
- Semaphore `mutex` initialized to the value 1
- Semaphore `full` initialized to the value 0
- Semaphore `empty` initialized to the value n

Bounded Buffer Problem (Cont.

- The structure of the producer process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```

Bounded Buffer Problem (Cont)

- ❓ The structure of the consumer process

```
Do {  
    wait(full);  
    wait(mutex);  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```


Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - Readers – only read the data set; they do **not** perform any updates
 - Writers – can both read and write
- Problem – allow multiple readers to read at the same time
 - Only one single writer can access the shared data at the same time
- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
 - Data set
 - Semaphore `rw_mutex` initialized to 1
 - Semaphore `mutex` initialized to 1
 - Integer `read_count` initialized to 0

Readers-Writers Problem (Cont.

- The structure of a writer process


```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```

Readers-Writers Problem (Cont.)

- The structure of a reader process

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
  
    /* ... reading is performed ... */  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```

Dining-Philosophers Problem



- Philosophers spend their lives alternating thinking and eating
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
 - Need both to eat, then release both when done
- In the case of 5 philosophers
 - Shared data
 - Bowl of rice (data set)
 - Semaphore **chopstick** [5] initialized to 1

Dining-Philosophers Problem Algorithm

- The structure of Philosopher *i*:

```
do {  
    wait (chopstick[i] );  
    wait (chopstick[ (i + 1) % 5] );  
  
    // eat  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
    // think  
  
} while (TRUE);
```

- What is the problem with this algorithm?

- Deadlock handling
 - Allow at most 4 philosophers to be sitting simultaneously at the table.
 - Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section).
 - Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.

Problems with Semaphores

- Incorrect use of semaphore operations:
 - signal (mutex) wait (mutex)
 - wait (mutex) ... wait (mutex)
 - Omitting of wait (mutex) or signal (mutex) (or both)
- Deadlock and starvation are possible.

Monitors


- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- *Abstract data type*, internal variables only accessible by code within the procedure
- Only one process may be active within the monitor at a time
- But not powerful enough to model some synchronization schemes

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { ... }


    procedure Pn (...) {.....}

    Initialization code (...) { ... }
}
}
```

Schematic view of a Monitor



Monitor with Condition Variable



End of Chapter 5

U.U.Samantha Rajapaksha

Senior Lecturer

Coordinator-M.Sc. in IT

Faculty of Computing

B.Sc.(Engineering) Moratuwa, M.Sc. IT (SLIIT)

Email: samantha.r@sliit.lk

Tel: 112301904 Ext: 4116

Web: www.sliit.lk