



SLIIT

Discover Your Future

IT2060/IE2061

Operating Systems and System Administration

Lecture 08

Memory Management

U. U. Samantha Rajapaksha

M.Sc.in IT, B.Sc.(Engineering) University of Moratuwa

Senior Lecturer SLIIT

Samantha.r@slit.lk



SLIIT
FACULTY OF COMPUTING

Chapter 8: Memory Management

- Background
- Swapping
- Contiguous Memory Allocation
- Paging

Main Memory – Chapter 8

- Memory is a large array of words or bytes, each with its own address
 - It is a repository of quickly accessible data shared by CPU and I/O devices.
 - Memory and registers are the only storage that CPU can directly access
- Main memory is a volatile storage device
 - It loses its contents in the case of system failure.
- A program must be mapped to absolute addresses and loaded into memory.
- Selection of a memory management scheme for a specific system depends on many factors, especially on the hardware design of the system.
- The OS is responsible for the following activities:
 - ❖ Keep track of which parts of memory are being used and by whom.
 - ❖ Decide which processes to load next when memory space becomes available.
 - ❖ Allocate and deallocate memory.



For n bits address, the memory capacity = 2^n bytes

Binary Address	Hex	Memory Bytes
0000 0000 0000 0000	0000	
0000 0000 0000 0001	0001	
0000 0000 0000 0010	0002	
0000 0000 0000 0011	0003	
0000 0000 0000 0100	0004	
0000 0000 0000 0101	0005	
		...
0000 0000 0100 1001	0049	
0000 0000 0100 1010	004A	
0000 0000 0100 1011	004B	
		...
1111 1111 1111 1111	FFFF	

Figure 1.2: Memory and Addresses

Background

- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are only storage CPU can access directly
- Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- Register access in one CPU clock (or less)
- Protection of memory required to ensure correct operation

Address Binding

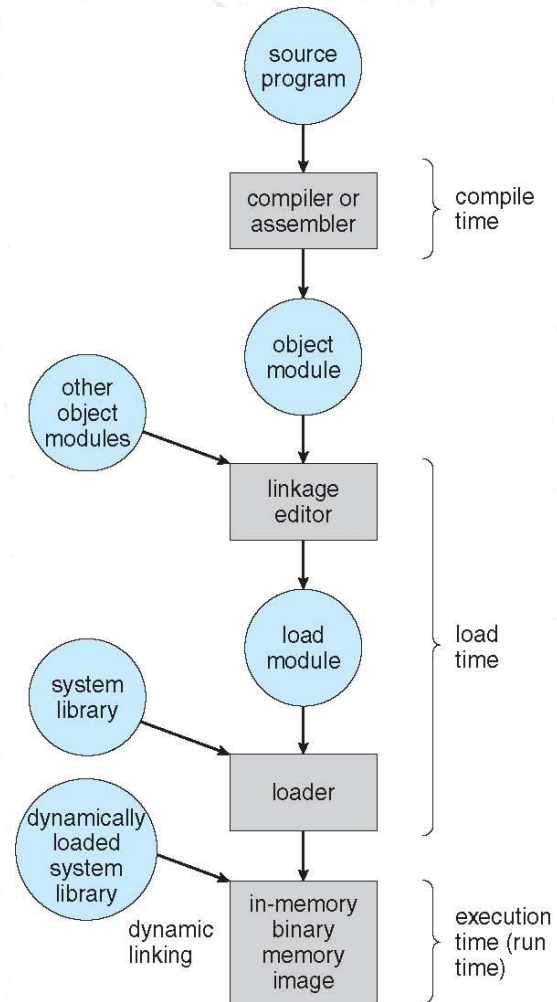
- Programs on disk, ready to be brought into memory to execute form an **input queue**
 - Without support, must be loaded into address 0000
- Inconvenient to have first user process physical address always at 0000
 - How can it not be?
- Further, addresses represented in different ways at different stages of a program's life
 - Source code addresses usually symbolic
 - Compiled code addresses **bind** to relocatable addresses
 - i.e. "14 bytes from beginning of this module"
 - Linker or loader will bind relocatable addresses to absolute addresses
 - i.e. 74014
 - Each binding maps one address space to another



Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages
 - **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
 - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
 - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - Need hardware support for address maps (e.g., base and limit registers)

Multistep Processing of a User Program



Logical vs. Physical Address Space


- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program



Swapping

- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
 - Total physical memory space of processes can exceed physical memory
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk

Schematic View of Swapping



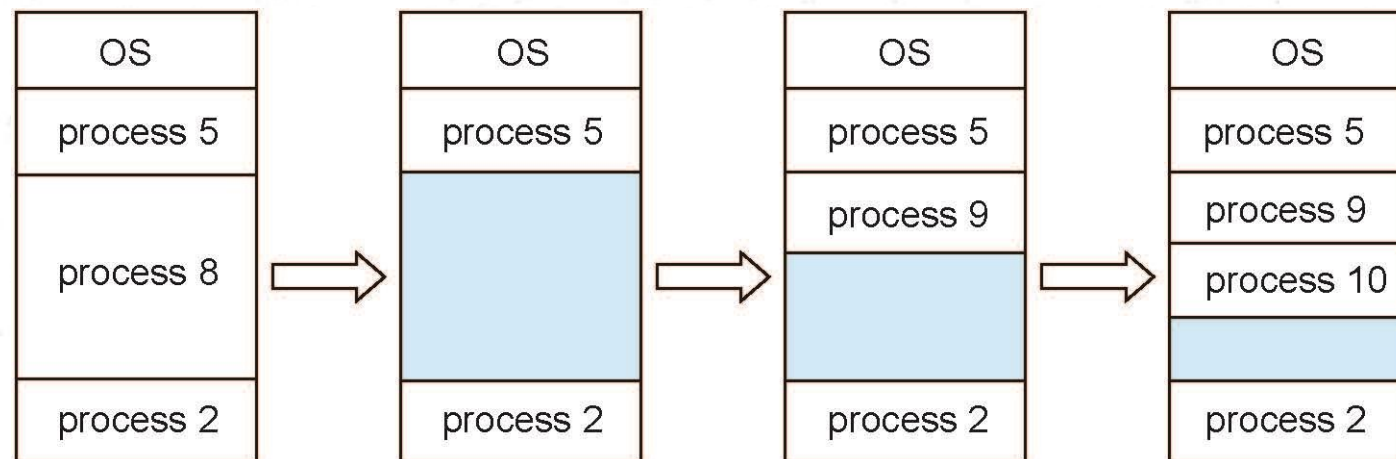
Contiguous Allocation

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually into two **partitions**:
 - Resident operating system, usually held in low memory with interrupt vector
 - User processes then held in high memory
 - Each process contained in single contiguous section of memory

Multiple-partition allocation

- Multiple-partition allocation

- Degree of multiprogramming limited by number of partitions
- **Variable-partition** sizes for efficiency (sized to a given process' needs)
- **Hole** – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Process exiting frees its partition, adjacent free partitions combined
- Operating system maintains information about:
 - a) allocated partitions b) free partitions (hole)



Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes?

- **First-fit**: Allocate the **first** hole that is big enough
- **Best-fit**: Allocate the **smallest** hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
- **Worst-fit**: Allocate the **largest** hole; must also search entire list
 - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization



Fragmentation

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - $1/3$ may be unusable -> **50-percent rule**



Fragmentation (Cont.)

- Reduce external fragmentation by **compaction**
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible *only* if relocation is dynamic, and is done at execution time
 - I/O problem
 - Latch job in memory while it is involved in I/O
 - Do I/O only into OS buffers
- Now consider that backing store has same fragmentation problems


Paging

- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size **N** pages, need to find **N** free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Backing store likewise split into pages
- Still have Internal fragmentation



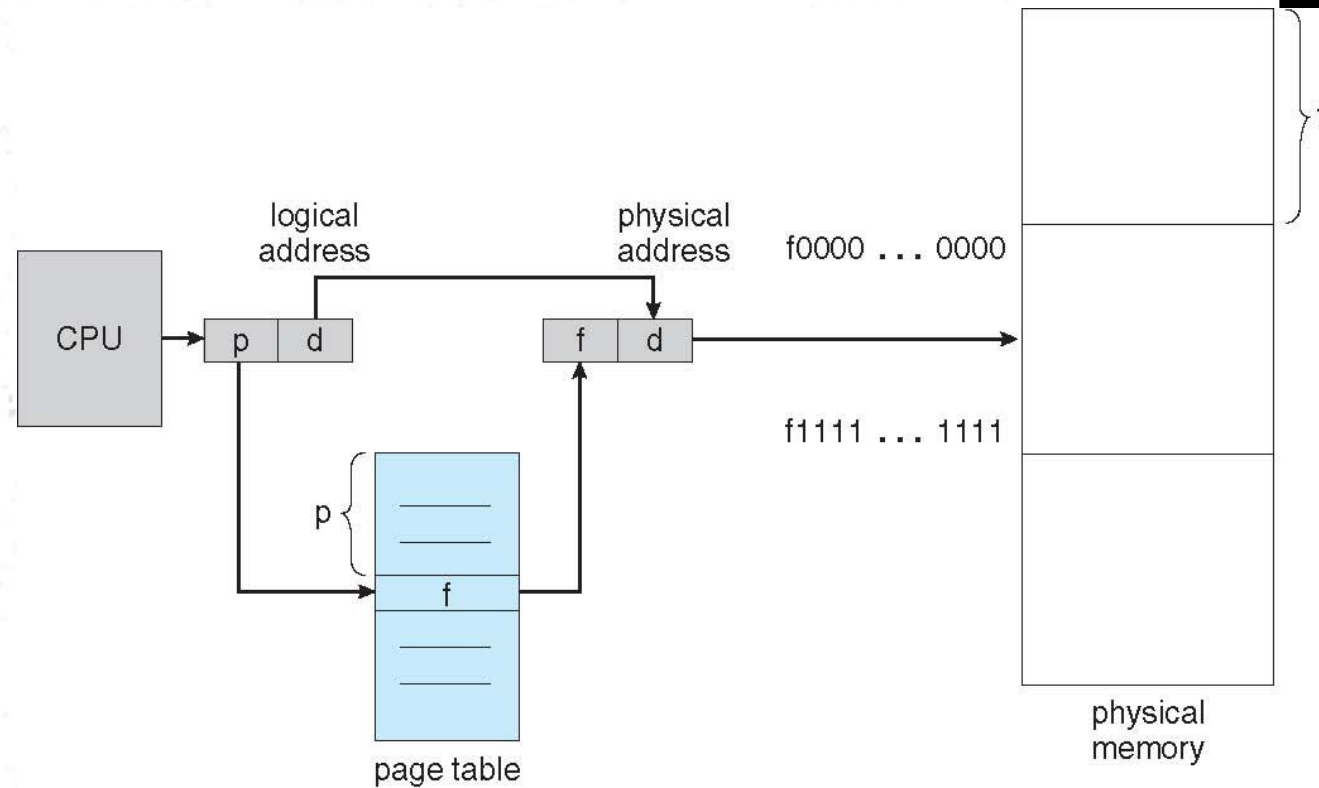
Address Translation Scheme

- Address generated by CPU is divided into
 - **Page number** (p) – used as an index into a **page table** which contains base address of each page in physical memory
 - **Page offset** (d) – combined with base address to define the physical memory address that is sent to the memory unit

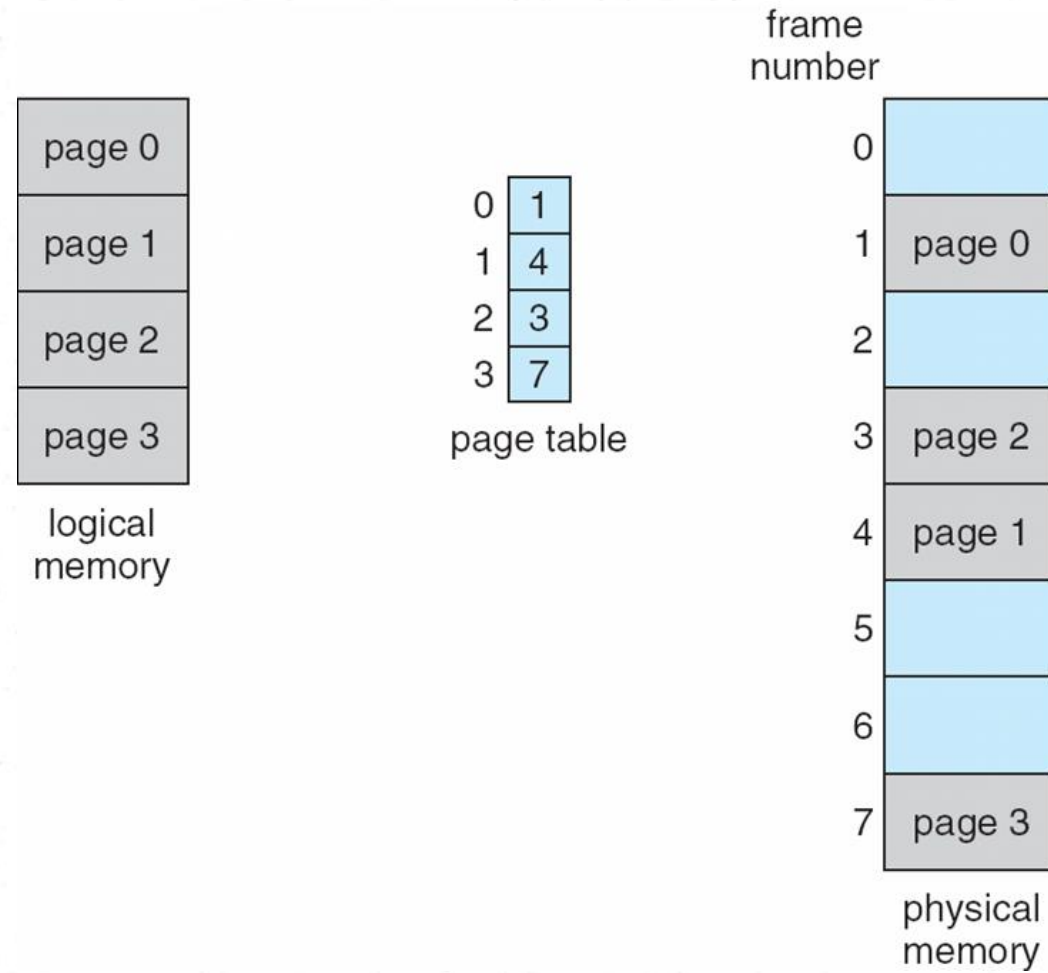


- For given logical address space 2^m and page size 2^n


Paging Hardware



Paging Model of Logical and Physical Memory



Paging Example





$n=2$ and $m=4$ 32-byte memory and 4-byte pages

Free Frames

free-frame list

14
13
18
20
15




(a)

Before allocation


free-frame list

15



0	14
1	13
2	18
3	20

new-process page table



(b)

After allocation


Associative Memory

- Associative memory – parallel search

Page #	Frame #

- Address translation (p, d)
 - If p is in associative register, get frame # out
 - Otherwise get frame # from page table in memory

Paging Hardware With TLB



Effective Access Time

- Associative Lookup = ε time unit
 - Can be < 10% of memory access time
- Hit ratio = α
 - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
- **Effective Access Time (EAT)**

$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$



- Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider more realistic hit ratio -> $\alpha = 99\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$

Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - “invalid” indicates that the page is not in the process’ logical address space
 - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel



Valid (v) or Invalid (i) Bit In A Page Table



Page Table Structure

- Hierarchical Paging.
- Hashed Page Tables.
- Inverted Page Tables.

Hierarchical Paging


- Most modern computer systems support a very large logical address, e.g., 2^{32} to 2^{64} → page table becomes very large.

Example: For a system with 32 bit logical address, and page size = 4KB (12 bits offset)

- page table = 1 Million entries ($32 - 12 = 20$ bits) → 4 MB of physical address spaces for page table; $2^{20} * 4$ bytes/entry
- One solution is to divide the page table into smaller pieces
 - For a two-level paging scheme, the page table itself is also paged.

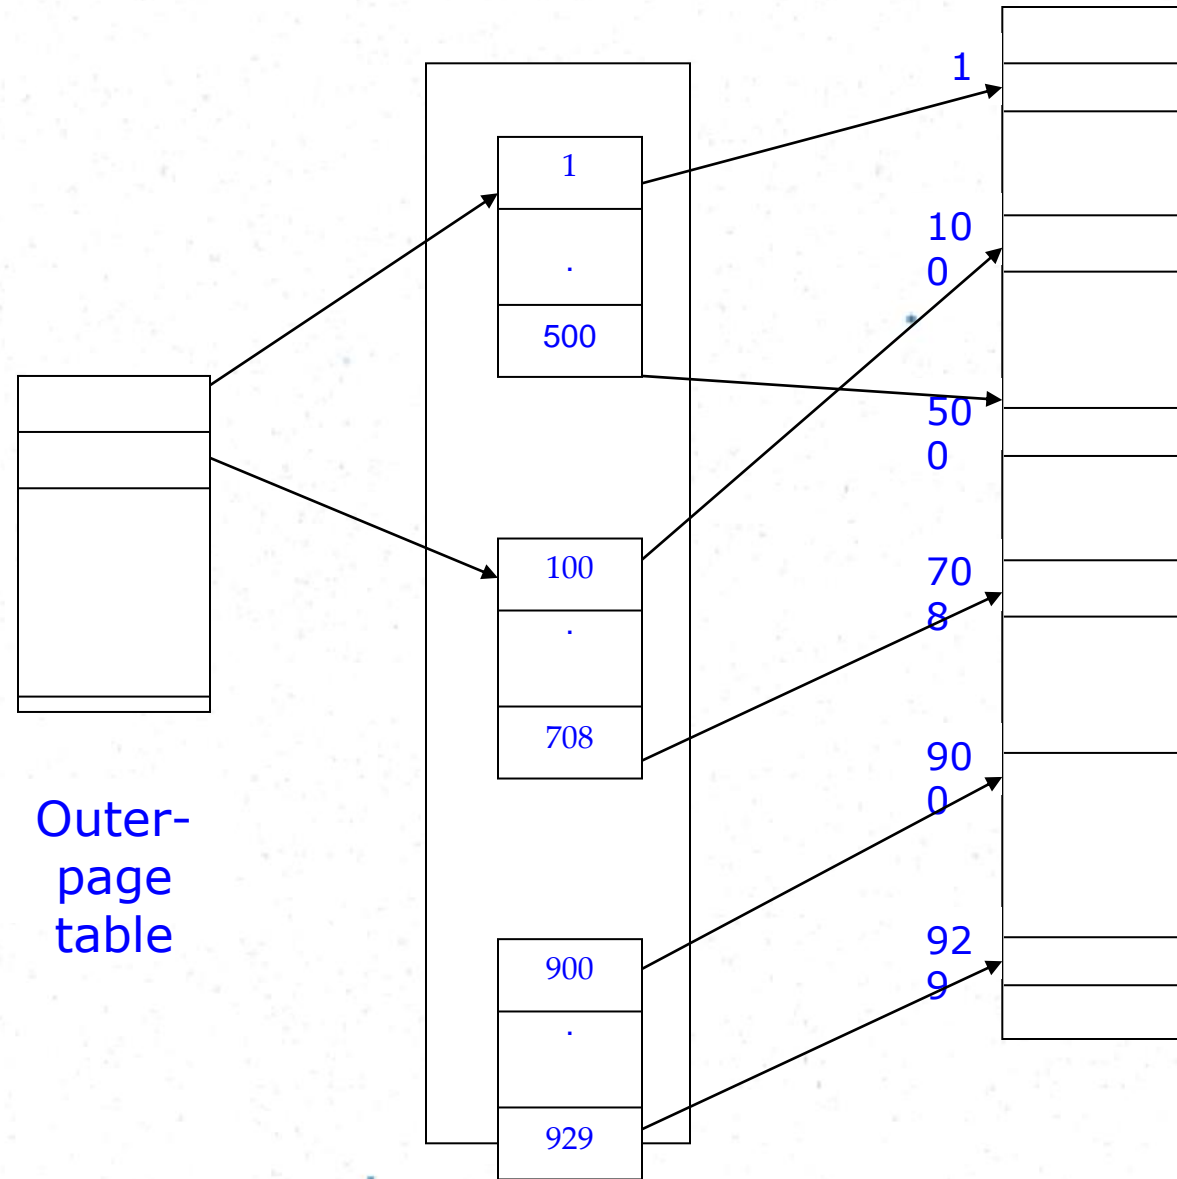
Two-level paging example

- A logical address (on 32-bit machine with 4K page size) is divided into:
 - a page number consisting of 20 bits.
 - a page offset consisting of 12 bits.
- Since the page table is paged, the page number is further divided into:
 - a 10 bit page number.
 - a 10 bit page offset.
- Thus, a logical address is as follows:



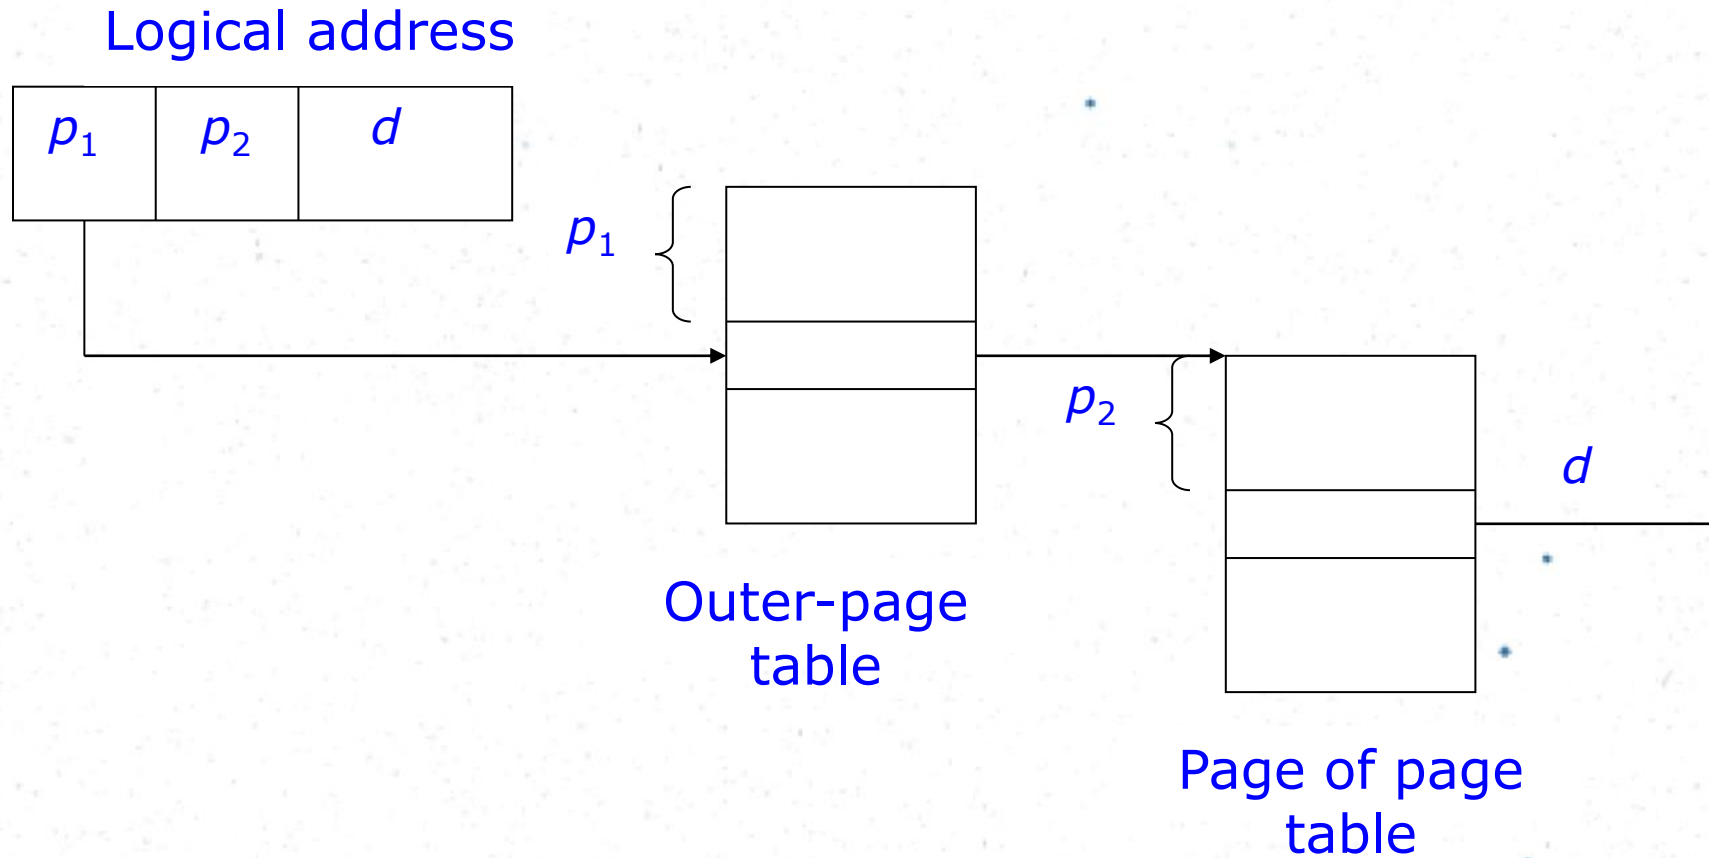
- * p_1 is an index into the *outer page table*, and p_2 is the displacement within the page of the *inner page table*
- * Translation starts from the outer page table

Two-level paging example



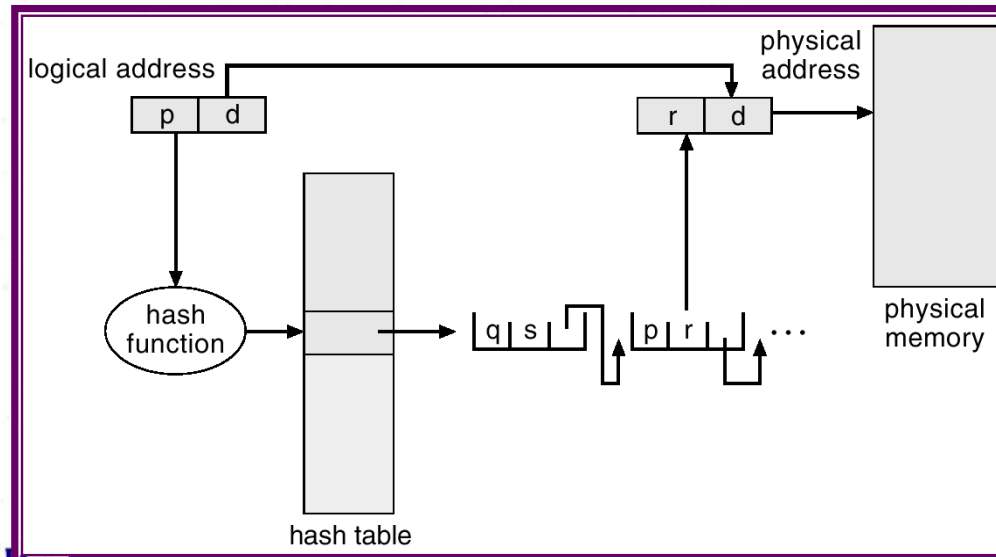
Address-translation scheme

Two-level 32-bit paging architecture



Hashed Page Tables

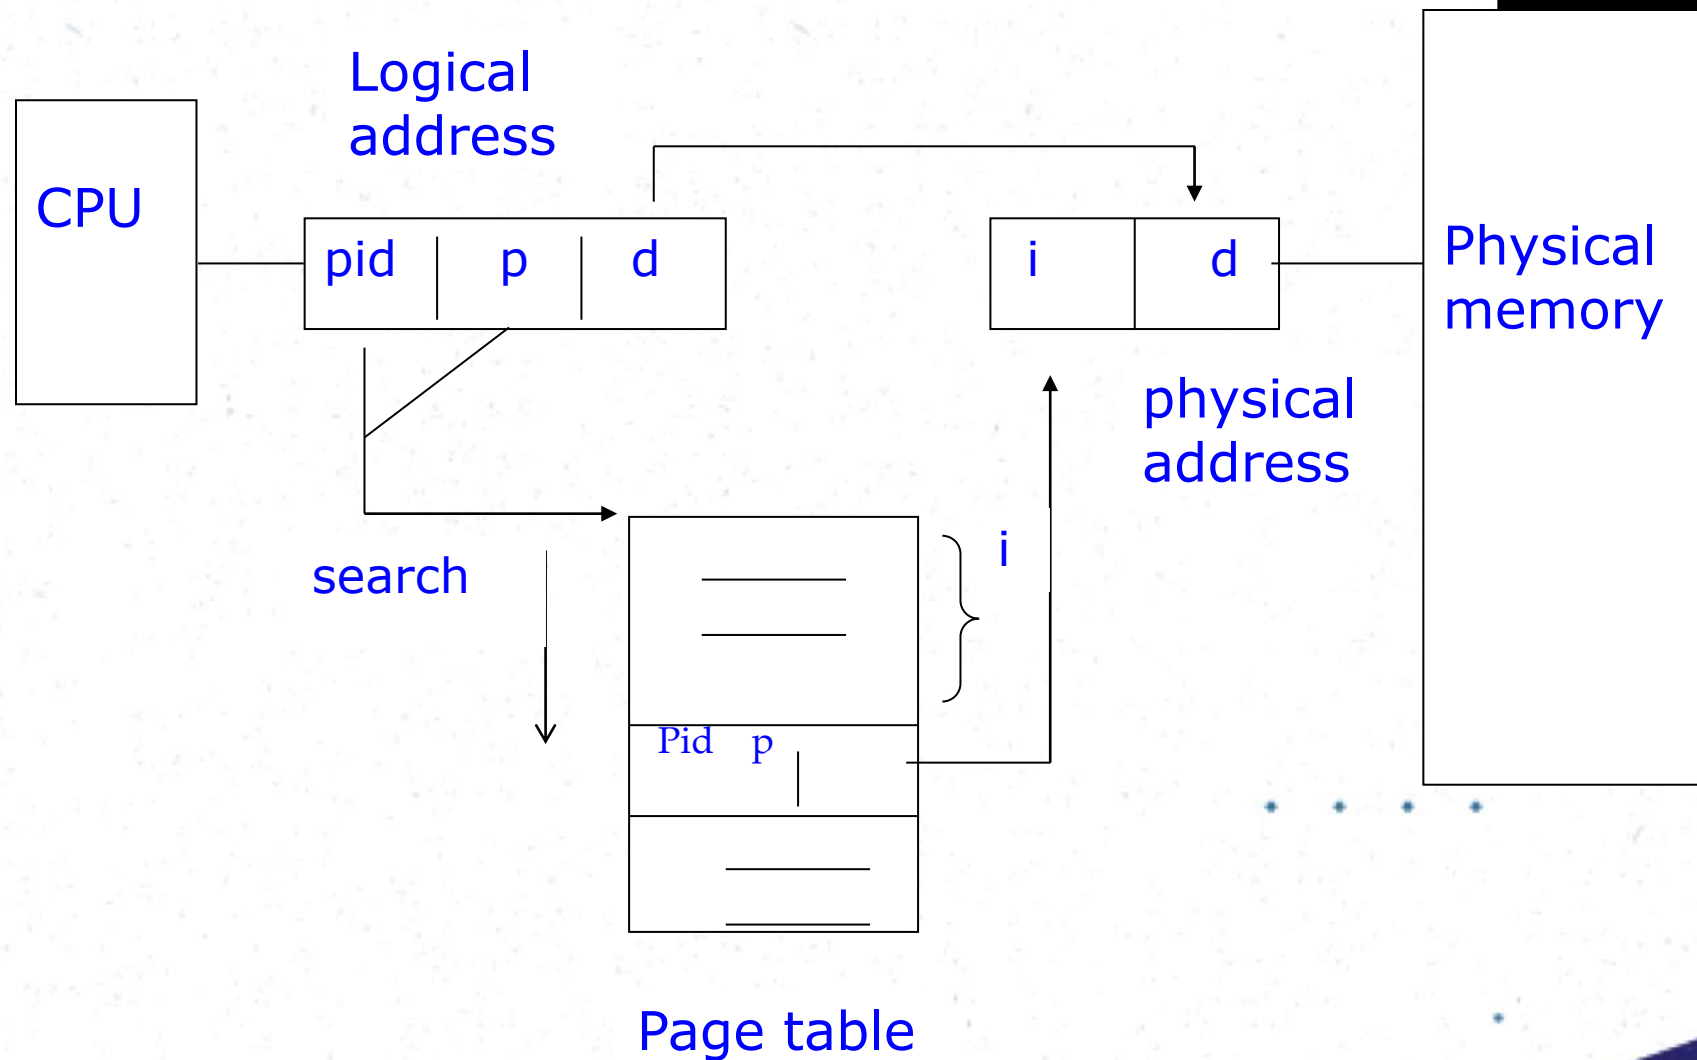
- Is commonly used when address spaces > 32 bits.
- The virtual page number is hashed into a page table.
 - This page table contains a chain of elements hashing to the same location.
 - Each element contains: virtual page number, the value of the mapped frame, pointer to the next element
- Virtual page number is compared in this chain searching for a match.
 - If a match is found, the corresponding physical frame is extracted.
- For address spaces > 64 bits, use Clustered page tables.
 - Similar to hash page table except each entry in page table refers to several pages (e.g., 16).



Inverted Page Table

- In the paging system, each process has a page table associated with it.
 - Drawback: each page table may consist of millions of entries → page tables consume large amount of physical memory.
- To solve this problem, use Inverted Page Table.
 - It uses ONLY one table that has one entry for each real page of memory.
- Each entry consists of the virtual address of the page stored in that real memory location, with information about the process (e.g., PID) that owns that page.
- This scheme *decreases* memory needed to store each page table, but *increases* time needed to search the table when a page reference occurs.
 - It may use hash table to limit the search to one – or at most a few page table entry.

Inverted page table architecture



Shared pages

- Other advantage of paging is the possibility of *sharing common* code (reentrant code also called pure code).
- Reentrant code is non-self-modifying code → its code never change during execution.
- Shared code.
 - One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).
- Private code and data.
 - Each process keeps a separate copy of the code and data.
 - The pages for the private code and data can appear anywhere in the logical address space.

Shared pages example

Process

ed 1
ed 2
ed 3
data 1

Page table
for p₁

3
4
6
1

Process

ed 1
ed 2
ed 3
data 2

3
4
6
7

Page table
for p₂

Process

ed 1
ed 2
ed 3
data 3

Page table
for p₃

3
4
6
2

0

1

2

3

4

5

6

7

8

9

10

data 1

data 3

ed 1

ed 2

ed 3


data 2

Segmentation


- Memory management with paging makes separation between user view of memory and the actual physical memory.
- Segmentation is memory-management scheme that supports user view of memory.
- From user's view, a program is a collection of segments, and a segment is a logical unit such as:
 - Main program.
 - Procedure / Function.
 - Local variables, global variables.
 - Common block.
 - Stack.
 - Symbol table, arrays.
- Each segment has a name and length; the addresses specify both the segment name and an offset.
- Program specifies each address by two quantities: a segment name, and an offset (within the segment).

Logical view of segmentation

User space
memory



Physical
memory



Segmentation architecture

- Logical address consists of a two tuple: $\langle \text{segment-number}, \text{offset} \rangle$.
- *Segment table* – maps two-dimensional user-defined addresses into one-dimensional physical addresses.
 - Each table entry has:
 - *Base* – contains the starting physical address where the segments reside in memory.
 - *Limit* – specifies the length of the segment.
- *Segment-table base register* (STBR) points to the segment table's location in memory.
- *Segment-table length register* (STLR) indicates the number of segments used by a program;
 - segment number s is legal if $s < \text{STLR}$.

Segmentation architecture (cont.)

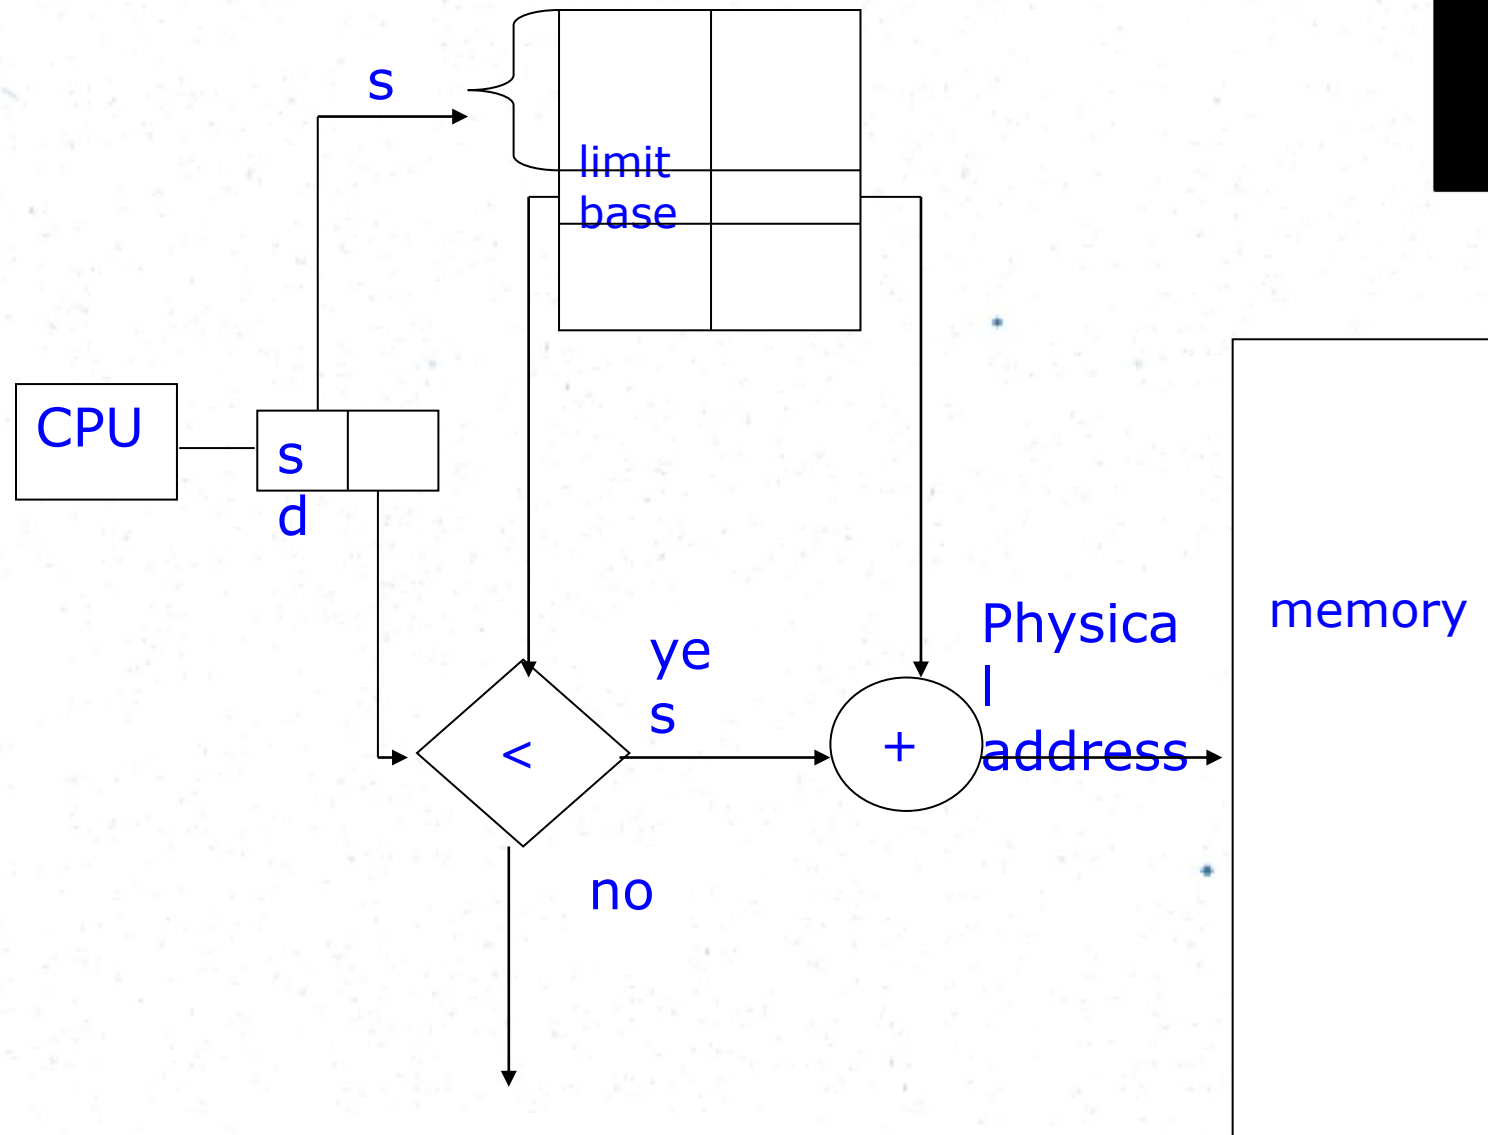
- Sharing.
 - Shared segments.
 - Same segment number.
- Allocation
 - First fit/best fit.
 - External fragmentation.

- Protection


With each entry in segment table associate:

- Validation bit = 0 \rightarrow illegal segment.
- Read/write/execute privileges.
- Protection bits associated with segments; code sharing occurs at segment level.
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem.

Segmentation Hardware





Trap; addressing
error



	limit	base
0	25286	43062
1	4425	68348

Segment table process P₁

4306
2
6834
8
7277
3
9000
3
855
3



	limit	base
0	25286	43062
1	8850	90003

Segment table process P₂